Your viewpoint exposure report

Topic: gender and relationships · last 30 days · MOCK-UP, example user

1 How much of your feed is this topic

46 of 2,200 videos (2.1%). Typical user: 1.6%.

Watching a lot of one topic is interest, not a cocoon. Part 2 is about what you see within it.

2 Within the topic: are you seeing one side more?

Compared with the same topic in the non-personalised feed.

Any opposing narrative

Not enough evidence yet

To tell a 2x difference from chance needs about 225 videos on this topic;

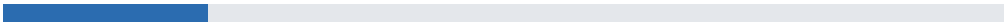

you have seen 46 (20.5% of what is needed).

Group labels

Not enough evidence yet

To tell a 2x difference from chance needs about 6,475 videos on this topic;

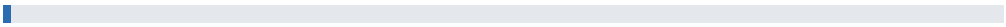

you have seen 46 (0.7% of what is needed).

We show a verdict only when the evidence is sufficient. For most users and rare narratives it never is.

3 What you can do, and what it may cost

Expected change in this narrative's share vs doing nothing. Simulation, upper scenario: assumes the recommender can tell the narrative apart; not measured on the app.

	effect	topic kept
Delete these videos from my history *lasts after you stop; needs accurate labels*	large	99%
Hide these videos (filter) *works only while on; keeps 86% with a weaker classifier*	large	95%
"Not interested", if it acts as negative feedback *depends on how the platform uses it*	medium	99%
"Not interested", if it behaves as in our OneRec test *barely helps; in OneRec it raised the topic's share*	little or none	100%
Turn off personalisation *removes only the personal layer*	small	99%
Skip these videos quickly	small	99%
Watch other videos on the topic to the end *ordinary videos barely offset it*	little or none	100%

Leave the topic entirely: the narrative disappears, and so does everything else on the topic.

4 What this report does not say

- It does not judge which view is right, or label any view as harmful.
- It does not infer what you believe; it describes only what you were shown.
- It does not claim the platform intends any pattern.
- Comparing with the non-personalised feed shows only your personal layer; a skew built into the platform's overall ranking would not appear here.
- Labels come from a language model and can be wrong.